SEER · FAO OPERATING PLAN CONTEST

FOREGUILD
THE AUTONOMOUS LABOR ECONOMY

Competition Edition · Revenue-First Operating Plan

A practical blueprint for turning AI-agent capability into trusted, repeatable, paid work—and using futarchy to improve measurable organizational decisions.

STRATEGIC THESIS	OPERATING PROMISE	GOVERNANCE PROMISE
Sell completed outcomes, not access to agents.	Every job has scope, evidence, acceptance and recourse.	Use markets for measurable choices; retain hard safety limits.

Prepared for the Seer FAO Operating Plan Contest

November 2026 · Independent proposal · Version 2.0

This is a proposed operating plan. It does not represent existing Foreguild traction, deployed contracts, revenue, partnerships, or approved governance.

Executive Decision Brief

Recommendation: launch Foreguild as a narrow, quality-controlled marketplace for repeatable digital work. Start with three categories—structured research, spreadsheet/data operations, and software QA/documentation—then expand only when repeat purchase and contribution margin are demonstrated.

The marketplace should solve the buyer's actual problem: getting a defined result delivered on time, with verifiable evidence and a clear remedy when quality falls short. AI agents are the supply-side engine; they are not the product promise. The product promise is dependable work.

Revenue should come first from a fee on accepted work, then from buyer subscriptions and enterprise/API services. Agent tooling can become a fourth revenue stream after Foreguild has enough demand to make paid provider tools valuable. Token appreciation, treasury speculation, and token issuance should not be required for operating viability.

Futarchy should govern bounded strategic choices with measurable outcomes—for example, which pricing package improves contribution margin without increasing disputes. It should not autonomously control emergency response, user safety, legal compliance, custody, or unrestricted treasury transfers.

The 12-month plan is gated, not calendar-driven: discover willingness to pay; prove paid completion; demonstrate repeat usage and positive variable contribution; then scale selected categories.

1. The Business in One Page

Element	Operating design
Customer	Small agencies, startup operations teams, research teams and crypto-native organizations with recurring digital work.
Job to be done	Turn a well-defined task into an accepted deliverable with predictable cost, timing and recourse.
Supply	Vetted AI agents and agent operators, routed by demonstrated capability and reliability.
Core workflow	Specify → fund → match → execute → verify → accept/dispute → settle → rate.
Primary revenue	Marketplace fee on successfully accepted task value.
Defensibility hypothesis	Accumulated task specifications, outcome-level reliability data, verification recipes, buyer workflow integrations and calibrated decision markets.
North-star metric	Monthly accepted task value from repeat buyers, accompanied by positive contribution per completed job.

2. Beachhead Market and Category Choice

The initial market should be selected by operational suitability, not by how impressive the demo looks. A good first category has frequent demand, clear deliverables, low-to-moderate downside if an agent fails, and inexpensive verification.

Category	Example deliverable	Verification	Launch stance
Research operations	Source-backed competitor brief or structured public-data research.	Required URLs, claim-source mapping, date checks, completeness rubric.	Pilot
Data operations	Cleaned CSV, normalized records, deduplication or CRM enrichment.	Schema checks, row counts, duplicate tests, sample review.	Pilot
Software QA/docs	Test report, reproducible bug ticket, documentation update.	Test execution, reproduction steps, lint/build checks, reviewer sign-off.	Pilot
High-stakes advice	Medical, legal, investment or safety-critical recommendations.	Requires qualified human responsibility and specialized controls.	Exclude initially

Why these three? They are digitally delivered, can be broken into bounded units, and allow objective checks or structured human review. The team should validate actual demand through interviews and paid pilots before committing engineering capacity.

3. Product Blueprint and Job Lifecycle

The job object is the core product primitive. Each job should contain: buyer, task category, scope, input references, output format, acceptance tests, deadline, price, escrow status, agent permissions, evidence bundle, review window and dispute state.

Lifecycle and control points

- **Intake:** convert a natural-language request into a structured brief; ask clarifying questions; reject prohibited or unbounded work.
- **Quote:** show fixed price or capped bid, estimated delivery, platform fee, verification method and refund/dispute terms before funding.
- **Fund:** segregate customer escrow from operating funds. Display the exact release conditions.
- **Route:** select agents using capability evidence, acceptance history, defect rate, on-time rate, price, capacity and task risk.
- **Execute:** provide only task-required data and permissions; sandbox untrusted inputs and keep secrets out of prompts.
- **Verify:** run deterministic tests where possible; require human review when quality is subjective or impact is material.
- **Accept or dispute:** buyer receives a concise evidence bundle and a time-limited review window; disputes pause only the contested funds where feasible.
- **Settle and learn:** release payment under published rules and update reputation from transaction outcomes.

Example transaction: a buyer pays $120 for a cleaned, deduplicated prospect CSV. The specification defines required columns, allowed sources, deduplication rules, sample size and delivery deadline. An agent returns the CSV plus a validation report. Automated checks test schema, nulls and duplicates; the buyer reviews a sample. On acceptance, escrow releases provider payment and the platform records the fee, cost-to-verify and outcome.

4. Revenue Architecture and Pricing Experiments

The first monetization principle is alignment: Foreguild earns more when buyers receive accepted value. Pricing must be simple enough to understand before a task starts and flexible enough to test across categories.

Stream	Initial offer (proposal)	Monetization trigger
Marketplace fee	Test 8%, 10% and 12% fee levels; initial planning case uses 10% of accepted task value.	Only after successful completion/acceptance.
Buyer Pro	Illustrative $49–$199/month tiers for saved workflows, team seats, usage analytics, priority routing and audit exports.	When a buyer repeats enough to value workflow and admin features.
Enterprise/API	Custom minimum commitment plus usage pricing for private pools, integrations, policy controls and reporting.	After a successful paid pilot and clear security requirements.
Agent tooling	Optional provider analytics, monitoring, scheduling and operational integrations.	Only after organic paid demand exists; avoid supply-side friction early.

Pricing experiment: recruit comparable buyer cohorts and rotate clearly disclosed fee offers or subscription packages. Compare conversion, repeat rate, gross profit per buyer and dispute rate. Do not optimize conversion alone; a lower fee that attracts low-quality or one-off demand may reduce long-term contribution.

5. Financial Model: Three 12-Month Scenarios

All figures below are illustrative planning assumptions, not actual Foreguild financials, a forecast, or verified market benchmarks. They are designed to make the operating logic auditable. USD-equivalent values are used; actual chain fees, settlement costs, taxes and asset volatility require validation.

Metric (12 months)	Conservative	Base	Upside
Accepted jobs	500	2,000	5,000

Metric (12 months)	Conservative	Base	Upside
Average task value	$80	$100	$120
Gross task value	$40,000	$200,000	$600,000
Marketplace take rate	8%	10%	10%
Marketplace revenue	$3,200	$20,000	$60,000
Assumed variable cost / job	$4.00	$3.00	$2.50
Variable costs total	$2,000	$6,000	$12,500
Contribution before fixed costs	$1,200	$14,000	$47,500

Formula: marketplace revenue = accepted jobs × average task value × take rate. Contribution before fixed costs = marketplace revenue – assumed variable costs. Variable cost includes payment/settlement allowance, automated verification and support allocation. The model excludes subscription/API revenue, payroll, engineering, legal, security audits, marketing, tax and financing costs.

Illustrative job-level economics

Per $100 accepted task	Amount	Treatment
Gross task value	$100.00	Customer-funded job value; not platform revenue.
10% platform fee	$10.00	Illustrative realized revenue after acceptance.
Settlement allowance	($1.00)	Assumption; validate per chain and payment route.
Verification/support	($2.00)	Assumption; track actual human minutes and tool cost.
Contribution	$7.00	Before fixed operating costs.

Decision rule: do not scale a task category unless contribution per accepted job is positive after realistic variable costs, repeat purchase is visible, and quality guardrails hold. If verification is expensive, either improve automation, raise price, narrow the task, or stop offering that category.

6. Competitive Positioning and Defensibility

Foreguild should not claim that it is automatically better than every freelance platform, agent framework or crypto marketplace. Its positioning should be tested against the alternatives buyers actually use today.

Alternative	Buyer trade-off to investigate	Foreguild response to prove
Human freelance marketplace	Human accountability and flexible work, but coordination, turnaround and repeatability may vary by task.	Faster standardized execution for bounded tasks, with clear verification and recourse.
Direct AI tools / agent frameworks	Low marginal execution cost, but buyer may still need to orchestrate, inspect and integrate results.	Outcome-oriented workflow: matching, escrow, evidence, acceptance and dispute handling.
Internal operations team	High control and context, but capacity and cost may limit flexible task throughput.	Elastic external capacity for well-specified overflow work; keep sensitive decisions internal.
Generic bounty boards	Useful for open calls, but task quality, evaluation and ongoing service levels can be inconsistent.	Structured task objects, repeatable workflows, verified agent reputation and payment rules.

Defensibility is a hypothesis to earn, not a feature checklist. The strongest potential moat is a proprietary dataset linking task specifications to verified outcomes: which agent performs which task, at what price, with what defect rate, review cost and buyer repeat behavior. Over time, this improves routing, pricing, verification and forecasting.

Proof required: publish cohort-level data showing Foreguild's cost, completion time, acceptance rate and repeat purchase against the buyer's current alternative. Do not publish unsupported superiority claims.

7. Futarchy: A Decision System with Guardrails

Futarchy is useful only when the organization can define a measurable outcome, a time horizon and a credible way to resolve the result. A market signal is an input to governance—not a magical truth oracle.

Proposed decision pipeline

Stage	Mechanism	Control
1. Proposal	Submit a policy option, rationale, owner, budget, metric and downside risks.	Reject vague proposals; require comparable alternatives.
2. Metric lock	Pre-register outcome, baseline, time window and data source.	No changing the metric after market positions are visible.
3. Forecast / market	Participants forecast the metric under each option; conditional markets may be used where legally and technically appropriate.	Minimum liquidity, participation and conflict disclosures.
4. Decision	Choose the option with the stronger expected outcome only if the pre-set margin and guardrails are met.	Quorum, time delay and documented review for material decisions.
5. Execution	Issue a bounded mandate with spend cap, owner, deadline and rollback condition.	No autonomous unrestricted treasury transfers.
6. Resolution	Resolve outcome from pre-agreed data and publish forecast-versus-actual report.	Independent resolver / appeal path for disputed data.

Pilot proposal: test whether a second verification step increases 30-day contribution margin per active buyer while keeping median completion time and dispute rate within fixed limits. Run the pilot in a low-risk category, cap spend, and compare results with a pre-registered baseline.

Hard limits outside market control

- Emergency security response, exploit containment, legal compliance, sanctions handling where applicable, user safety and privacy incident response.
- Custody rules, contract upgrade authority, constitutional rights and changes to the definition of accepted work.
- Treasury solvency, customer escrow segregation, per-agent spending caps and any transfer outside approved allowlists.
- High-stakes or irreversible decisions without accountable human review.

Market integrity controls: liquidity threshold; capped position or influence concentration; disclosure of conflicts; independent resolution source; cooling-off period for high-impact proposals; anti-manipulation monitoring; and a narrow, documented safety override with post-event review.

8. Agent Architecture, Safety and Human Oversight

- **Intake agent:** creates structured job briefs, clarifies ambiguity, classifies risk and blocks unsupported or disallowed tasks.
- **Routing agent:** scores eligible providers using verified capability, recent reliability, price, capacity and risk tier.
- **Execution agents:** work in isolated sandboxes with scoped tools and task-specific access; untrusted documents are treated as data, not instructions.
- **Verifier agent:** executes deterministic tests, checks source traceability, schema compliance and output completeness.
- **Human reviewer:** handles subjective acceptance, high-value work, appeals and risk escalations.
- **Finance agent:** reconciles balances and reports anomalies but cannot independently change treasury permissions or send unrestricted funds.

Agent reputation should be transaction-based: first-pass acceptance, on-time delivery, defect/rework rate, dispute outcomes and verified buyer feedback. New agents start with low job-value limits. Increase limits gradually after a statistically meaningful history; do not let self-attested skills alone unlock high-value work.

9. Treasury, Escrow and Dispute Operations

- **Separate ledgers:** customer escrow, platform revenue, operating treasury and reserves must be tracked separately. Escrow is a liability, not revenue.
- **Release logic:** publish acceptance window, automatic release conditions, partial acceptance rules and refund paths before funding.

- **Security:** use audited contracts where applicable, capped exposure, multisignature approvals, allowlisted destinations, transaction simulation and emergency pause/recovery plans.
- **Disputes:** preserve evidence, distinguish objective failure from subjective dissatisfaction, pause only the contested amount where possible, and provide an appeal path.
- **Reporting:** reconcile gross task value, provider payouts, fees, refunds, disputes and treasury movement in a regular public report.

10. Go-to-Market: Demand Before Supply

- **Design partners:** recruit 10–15 small agencies, startup operations teams and crypto-native organizations with recurring digital tasks.
- **Founder-led sales:** sell a named outcome (for example, “validated competitor brief by tomorrow”), not a generic AI-agent subscription.
- **Template-led acquisition:** publish sample tasks with input requirements, output examples, acceptance criteria and indicative pricing.
- **Partner channels:** developer communities, startup accelerators, DAO contributor networks and service agencies.
- **Referral loop:** reward only after a referred buyer completes and pays for a first job.
- **Supply curation:** onboard a small set of agents with demonstrated samples; expand supply only when buyer demand exists.

Primary growth metrics: paid buyer conversion, time to first accepted job, 30/90-day repeat buyer rate, gross profit per buyer and contribution per task. Secondary metrics include active qualified agents and fill rate. Vanity metrics do not determine scale decisions.

11. Twelve-Month Roadmap and Stage Gates

Phase	Timing	Deliverables	Gate to advance
Discover	M1–2	30 buyer interviews; 3 categories; manual concierge workflow; price tests.	10 buyers willing to pay; 20 scoped pilot tasks.
Paid pilot	M3–4	Escrow/controlled settlement, job templates, verification and dispute flow.	100 completed jobs; ≥70% accepted without dispute; measured positive variable contribution in at least one category.
Repeatability	M5–8	Buyer Pro test, routing, reputation, workflow integrations.	500 cumulative jobs; ≥30% repeat buyers; positive contribution in core category.
Selective scale	M9–12	Enterprise/API pilots; expand only proven categories; first futarchy pilot.	2,000 cumulative jobs; ≥40% repeat buyers; dispute rate <5% target.

These are proposed targets, not guaranteed outcomes. If a gate is missed, pause expansion and diagnose demand, quality, pricing or cost-to-serve before adding features or categories.

12. First 30 Days: Action Plan

Week	Actions	Evidence produced
1	Interview 15 buyers and 10 agent operators; collect real task examples, current alternatives, costs and willingness to pay.	Interview notes, task taxonomy, ranked pain points.
2	Choose 3 templates; define acceptance criteria, data boundaries, price and dispute rules; recruit pilot agents.	Pilot-ready job specs and provider scorecards.
3	Run paid jobs with controlled escrow/settlement; measure reviewer time, variable cost and defects.	Transaction-level unit economics and quality logs.
4	Analyze cohorts; remove unprofitable task types; revise pricing; automate the largest repeated bottleneck.	Go/no-go memo and next 60-day priorities.

13. KPI Dashboard and Decision Cadence

Dimension	Weekly / monthly measures	Management response
Demand	Qualified leads, paid conversion, first-job time, repeat rate, CAC.	Fix targeting and onboarding before increasing agent supply.
Marketplace	Fill rate, completion, acceptance, cancellation, dispute rate, delivery time.	Improve specifications, routing and verification; suspend weak categories.
Economics	Gross task value, realized fee revenue, contribution/job, variable cost, refund rate, runway.	Adjust price, workflow or category mix; do not confuse escrow with revenue.
Agent quality	On-time rate, first-pass acceptance, defects, rework, category-level performance.	Change routing weights, job limits or provider eligibility.
Governance	Market liquidity, participation, forecast calibration, decision accuracy, overrides.	Improve market design; suspend decisions when thresholds fail.

14. Key Risks and Mitigations

Risk	Mitigation
Fabricated or low-quality work	Task-specific tests, source traceability, evidence bundles, sampling and human review.
Fake agents / Sybil reputation	Risk-proportionate identity checks, transaction-based reputation, rate limits and anomaly detection.
Contract or treasury exploit	Audits, limited contract scope, capped exposure, multisig, emergency pause and recovery drills.
Prompt injection / data leakage	Sandboxing, least privilege, secret redaction, untrusted-input isolation and retention controls.
Thin or manipulated futarchy markets	Liquidity/participation thresholds, disclosures, concentration caps, independent resolution and safety override.
Regulatory / tax uncertainty	Jurisdiction-specific legal and tax review; transparent fee and risk disclosures.
Weak repeat demand	Validate with paid pilots; stop categories with poor retention or negative contribution.

15. Decisions Requested from the Founding Team

- Approve the initial three low-risk categories and explicit exclusions for the pilot.
- Select the fee experiment and the first buyer segment.
- Approve job acceptance, escrow, refund, dispute and appeal policies.
- Define operating runway, reserve policy and treasury approval limits.
- Choose one low-risk futarchy pilot with a pre-registered metric, time window and guardrails.
- Commit to a monthly transparency report covering jobs, acceptance, disputes, fees, variable costs and treasury movement.

Conclusion: Build Trust, Then Scale

Foreguild's opportunity is to make AI labor purchasable with the confidence buyers expect from a reliable service provider. That requires more than agent orchestration: it requires precise job definitions, verification, accountable settlement, trustworthy reputation and repeatable customer acquisition. The most credible path is to start narrow, prove paid outcomes, improve unit economics, and use futarchy only where measurable forecasts can add decision value.

North star: grow accepted work from repeat buyers while maintaining positive contribution per job, low disputes and strict control over customer funds and organizational risk.

Appendix A — Assumptions and Model Notes

- All scenario figures are illustrative, not actual Foreguild performance or a market forecast.

- Gross task value is the total value of buyer-funded jobs; it is not platform revenue.
- Marketplace revenue is modeled as accepted task value multiplied by take rate.
- Variable cost per job is an assumed blended allowance for settlement, verification and support; validate using real transaction logs.
- Fixed costs, payroll, legal, audits, marketing, tax, financing, subscription revenue and enterprise/API revenue are excluded from the scenario table.
- No token price, treasury yield, speculative asset appreciation or uncommitted grant is included in the operating case.

Appendix B — Validation Checklist Before Launch

- Buyer interviews confirm recurring pain and willingness to pay.
- Each pilot task has a testable definition of done and documented exclusions.
- Escrow and provider payout flows are tested; customer funds are segregated.
- Data access, retention, sandbox and incident procedures are documented.
- Dispute, refund, appeal and human escalation paths are staffed.
- Metrics are instrumented at job level and finance reconciles fees, payouts and refunds.
- Futarchy pilot metric, resolution source, liquidity threshold and safety override are pre-registered.

Appendix C — Brief and Source

Primary contest reference: Poidh, Bounty #1531, “FAO Operating Plan Contest,” https://poidh.xyz/base/bounty/1531. The contest brief is the basis for the requested Foreguild operating-plan framing. The proposal’s operating recommendations, pricing tests, scenario values and milestones are original assumptions for evaluation—not facts asserted by the contest organizer.

Important limitation: before submission, re-open the live bounty page and confirm its current status, exact deliverable instructions, submission method, eligibility, deadline and any newly added clarifications. This document does not itself submit an entry.